

```
In [14]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt,seaborn as sns
```

```
In [15]: train_df = pd.read_csv(r"C:\Users\sudheer\Downloads\Mobile_Price_Classification_train (1).csv")
train_df
```

Out[15]:

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	...	px_height
0	842	0	2.2	0	1	0	7	0.6	188	2	...	20
1	1021	1	0.5	1	0	1	53	0.7	136	3	...	905
2	563	1	0.5	1	2	1	41	0.9	145	5	...	1263
3	615	1	2.5	0	0	0	10	0.8	131	6	...	1216
4	1821	1	1.2	0	13	1	44	0.6	141	2	...	1208
...	...	...	...	...	...	...	...	...	...	...	...	...
1995	794	1	0.5	1	0	1	2	0.8	106	6	...	1222
1996	1965	1	2.6	1	0	0	39	0.2	187	4	...	915
1997	1911	0	0.9	1	1	1	36	0.7	108	8	...	868
1998	1512	0	0.9	0	4	1	46	0.1	145	5	...	336
1999	510	1	2.0	1	5	1	45	0.9	168	6	...	483

2000 rows × 21 columns



```
In [16]: test_df = pd.read_csv(r"C:\Users\sudheer\Downloads\Mobile_Price_Classification_test (1).csv")
test_df
```

Out[16]:

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	...	pc	px_height
0	1	1043	1	1.8	1	14	0	5	0.1	193	...	16	226
1	2	841	1	0.5	1	4	1	61	0.8	191	...	12	746
2	3	1807	1	2.8	0	1	0	27	0.9	186	...	4	1270
3	4	1546	0	0.5	1	18	1	25	0.5	96	...	20	295
4	5	1434	0	1.4	0	11	1	49	0.5	108	...	18	749
...	...	...	...	...	...	...	...	...	...	...	...	...	...
995	996	1700	1	1.9	0	0	1	54	0.5	170	...	17	644
996	997	609	0	1.8	1	0	0	13	0.9	186	...	2	1152
997	998	1185	0	1.4	0	1	1	8	0.5	80	...	12	477
998	999	1533	1	0.5	1	0	0	50	0.4	171	...	12	38
999	1000	1270	1	0.5	0	4	1	35	0.1	140	...	19	457

1000 rows × 21 columns



```
In [17]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   battery_power    2000 non-null   int64
1   blue             2000 non-null   int64
2   clock_speed      2000 non-null   float64
3   dual_sim         2000 non-null   int64
4   fc               2000 non-null   int64
5   four_g           2000 non-null   int64
6   int_memory       2000 non-null   int64
7   m_dep            2000 non-null   float64
8   mobile_wt        2000 non-null   int64
9   n_cores          2000 non-null   int64
10  pc               2000 non-null   int64
11  px_height        2000 non-null   int64
12  px_width         2000 non-null   int64
13  ram              2000 non-null   int64
14  sc_h             2000 non-null   int64
15  sc_w             2000 non-null   int64
16  talk_time        2000 non-null   int64
17  three_g          2000 non-null   int64
18  touch_screen     2000 non-null   int64
19  wifi             2000 non-null   int64
20  price_range      2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.2 KB
```

```
In [18]: test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id              1000 non-null   int64
1   battery_power    1000 non-null   int64
2   blue            1000 non-null   int64
3   clock_speed      1000 non-null   float64
4   dual_sim         1000 non-null   int64
5   fc               1000 non-null   int64
6   four_g           1000 non-null   int64
7   int_memory       1000 non-null   int64
8   m_dep            1000 non-null   float64
9   mobile_wt        1000 non-null   int64
10  n_cores          1000 non-null   int64
11  pc               1000 non-null   int64
12  px_height        1000 non-null   int64
13  px_width         1000 non-null   int64
14  ram              1000 non-null   int64
15  sc_h             1000 non-null   int64
16  sc_w             1000 non-null   int64
17  talk_time        1000 non-null   int64
18  three_g          1000 non-null   int64
19  touch_screen     1000 non-null   int64
20  wifi             1000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 164.2 KB
```

```
In [19]: x=train_df.drop('wifi',axis=1)
y=train_df['wifi']
```

```
In [20]: x=test_df.drop('wifi',axis=1)
y=test_df['wifi']
```

```
In [21]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,train_size=0.7,random_state=42)
x_train.shape,x_test.shape
```

```
Out[21]: ((700, 20), (300, 20))
```

```
In [22]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[22]: ▾ RandomForestClassifier
RandomForestClassifier()
```

```
In [23]: rf = RandomForestClassifier()
```

```
In [24]: params={'max_depth':[2,3,5,10,20],
'min_samples_leaf':[5,10,20,50,100,200],
'n_estimators':[10,25,30,50,100,200]}
```

```
In [25]: from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring='accuracy')
grid_search.fit(x_train,y_train)
```

```
Out[25]: ▸ GridSearchCV
▸ estimator: RandomForestClassifier
▸ RandomForestClassifier
```

```
In [26]: grid_search.best_score_
```

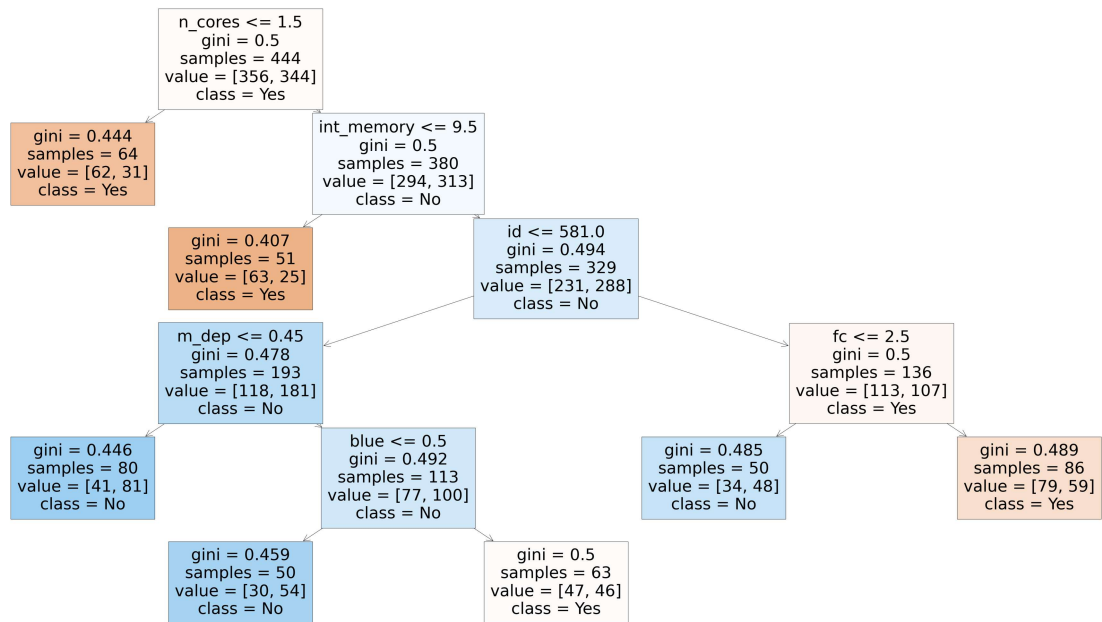
```
Out[26]: 0.5514285714285714
```

```
In [27]: rf_best=grid_search.best_estimator_
rf_best
```

```
Out[27]: ▾ RandomForestClassifier
RandomForestClassifier(max_depth=20, min_samples_leaf=50)
```

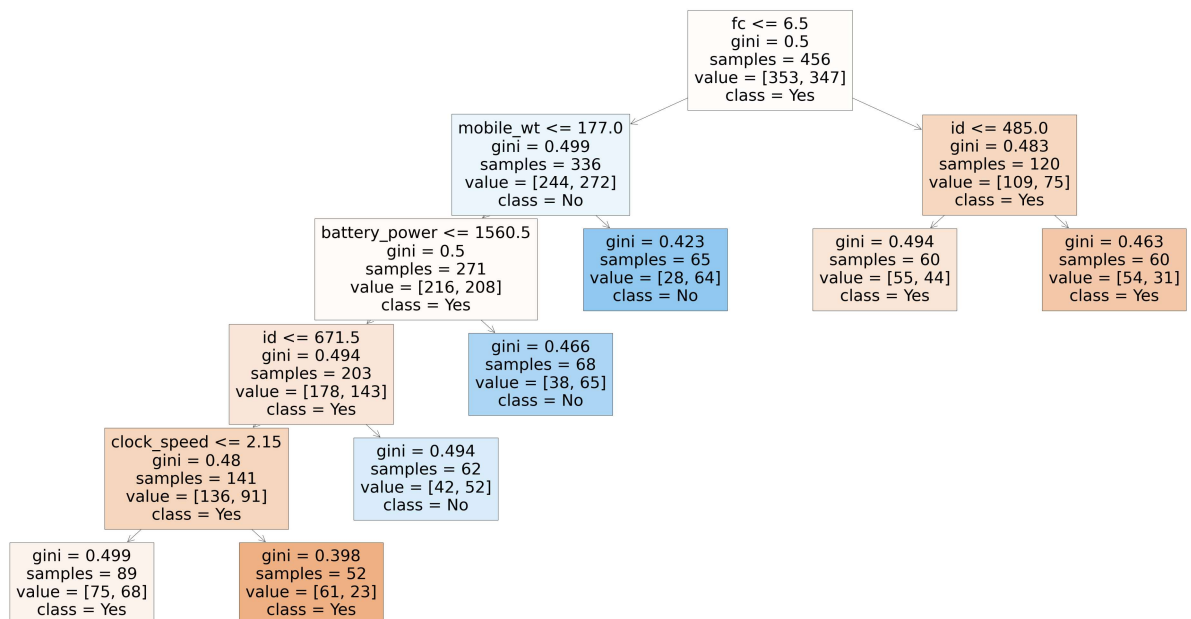
```
In [28]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[5], feature_names = x.columns,class_names=['Yes','No'],filled=True)
```

```
Out[28]: [Text(0.25, 0.9166666666666666, 'n_cores <= 1.5\ngini = 0.5\nsamples = 444\nvalue = [356, 344]\n\nclass = Yes'),
Text(0.125, 0.75, 'gini = 0.444\nsamples = 64\nvalue = [62, 31]\n\nclass = Yes'),
Text(0.375, 0.75, 'int_memory <= 9.5\ngini = 0.5\nsamples = 380\nvalue = [294, 313]\n\nclass = No'),
Text(0.25, 0.5833333333333333, 'gini = 0.407\nsamples = 51\nvalue = [63, 25]\n\nclass = Yes'),
Text(0.5, 0.5833333333333333, 'id <= 581.0\ngini = 0.494\nsamples = 329\nvalue = [231, 288]\n\nclass = No'),
Text(0.25, 0.4166666666666667, 'm_dep <= 0.45\ngini = 0.478\nsamples = 193\nvalue = [118, 181]\n\nclass = No'),
Text(0.125, 0.25, 'gini = 0.446\nsamples = 80\nvalue = [41, 81]\n\nclass = No'),
Text(0.375, 0.25, 'blue <= 0.5\ngini = 0.492\nsamples = 113\nvalue = [77, 100]\n\nclass = No'),
Text(0.25, 0.08333333333333333, 'gini = 0.459\nsamples = 50\nvalue = [30, 54]\n\nclass = No'),
Text(0.5, 0.08333333333333333, 'gini = 0.5\nsamples = 63\nvalue = [47, 46]\n\nclass = Yes'),
Text(0.75, 0.4166666666666667, 'fc <= 2.5\ngini = 0.5\nsamples = 136\nvalue = [113, 107]\n\nclass = Yes'),
Text(0.625, 0.25, 'gini = 0.485\nsamples = 50\nvalue = [34, 48]\n\nclass = No'),
Text(0.875, 0.25, 'gini = 0.489\nsamples = 86\nvalue = [79, 59]\n\nclass = Yes')]
```



```
In [29]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7], feature_names = x.columns,class_names=['Yes','No'],filled=True)
```

```
Out[29]: [Text(0.6363636363636364, 0.9166666666666666, 'fc <= 6.5\ngini = 0.5\nsamples = 456\nvalue = [353, 347]\nnclass = Yes'),
Text(0.45454545454545453, 0.75, 'mobile_wt <= 177.0\ngini = 0.499\nsamples = 336\nvalue = [244, 272]\nnclass = No'),
Text(0.36363636363636365, 0.58333333333333334, 'battery_power <= 1560.5\ngini = 0.5\nsamples = 271\nvalue = [216, 208]\nnclass = Yes'),
Text(0.2727272727272727, 0.4166666666666667, 'id <= 671.5\ngini = 0.494\nsamples = 203\nvalue = [178, 143]\nnclass = Yes'),
Text(0.18181818181818182, 0.25, 'clock_speed <= 2.15\ngini = 0.48\nsamples = 141\nvalue = [136, 91]\nnclass = Yes'),
Text(0.09090909090909091, 0.08333333333333333, 'gini = 0.499\nsamples = 89\nvalue = [75, 68]\nnclass = Yes'),
Text(0.2727272727272727, 0.08333333333333333, 'gini = 0.398\nsamples = 52\nvalue = [61, 23]\nnclass = Yes'),
Text(0.36363636363636365, 0.25, 'gini = 0.494\nsamples = 62\nvalue = [42, 52]\nnclass = No'),
Text(0.45454545454545453, 0.4166666666666667, 'gini = 0.466\nsamples = 68\nvalue = [38, 65]\nnclass = No'),
Text(0.5454545454545454, 0.58333333333333334, 'gini = 0.423\nsamples = 65\nvalue = [28, 64]\nnclass = No'),
Text(0.8181818181818182, 0.75, 'id <= 485.0\ngini = 0.483\nsamples = 120\nvalue = [109, 75]\nnclass = Yes'),
Text(0.7272727272727273, 0.58333333333333334, 'gini = 0.494\nsamples = 60\nvalue = [55, 44]\nnclass = Yes'),
Text(0.9090909090909091, 0.58333333333333334, 'gini = 0.463\nsamples = 60\nvalue = [54, 31]\nnclass = Yes')]
```



```
In [30]: rf_best.feature_importances_
```

```
Out[30]: array([0.06031287, 0.07588404, 0.01249582, 0.08418056, 0.01220114,
0.05887866, 0.02418764, 0.07420338, 0.06991117, 0.12339559,
0.01764515, 0.04971102, 0.06579852, 0.10809968, 0.06200155,
0.02275696, 0.02744152, 0.04592882, 0.00219524, 0.00277066])
```

```
In [31]: imp_df=pd.DataFrame({'Varname':x_train.columns,'Imp':rf_best.feature_importances_})
imp_df.sort_values(by='Imp',ascending=False)
```

Out[31]:

	Varname	Imp
9	mobile_wt	0.123396
13	px_width	0.108100
3	clock_speed	0.084181
1	battery_power	0.075884
7	int_memory	0.074203
8	m_dep	0.069911
12	px_height	0.065799
14	ram	0.062002
0	id	0.060313
5	fc	0.058879
11	pc	0.049711
17	talk_time	0.045929
16	sc_w	0.027442
6	four_g	0.024188
15	sc_h	0.022757
10	n_cores	0.017645
2	blue	0.012496
4	dual_sim	0.012201
19	touch_screen	0.002771
18	three_g	0.002195

In [ ]: